


Recap

24 December 2024 17:37



Steps for GPU Training

24 December 2024 17:37

1. Check GPU Availability

Before starting, verify if a GPU is available. If available, select it; otherwise, use the CPU.

```
python Copy code  
  
import torch  
  
# Check for GPU  
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')  
print(f"Using device: {device}")
```

2. Move the Model to GPU

Move your model to the selected device (`cuda` for GPU or `cpu`) so that all computations occur on the same device.

```
python Copy code  
  
model = MyModel() # Replace with your model  
model = model.to(device) # Move model to GPU
```

3. Modify the Training Loop by Moving Data to GPU

Ensure that each batch of data (features and labels) is moved to the GPU before processing. This ensures that both the model and data are on the same device.

```
for batch_features, batch_labels in train_loader:  
    # Move data to GPU  
    batch_features, batch_labels = batch_features.to(device), batch_labels.to(device)
```

4. Modify the Evaluation Loop by Moving Data to GPU

Similarly, ensure test data is moved to the GPU during evaluation. Disable gradient calculations using `torch.no_grad()` for efficiency.

```
with torch.no_grad():  
    for batch_features, batch_labels in test_loader:  
        # Move data to GPU  
        batch_features, batch_labels = batch_features.to(device), batch_labels.to(device)
```

5. Optimize the GPU Usage

To make the best use of GPU resources, apply the following optimizations:

a. Use Larger Batch Sizes

Larger batch sizes can better utilize GPU memory and reduce computation time per epoch (if memory allows).

```
python Copy code  
  
train_loader = DataLoader(train_dataset, batch_size=128, shuffle=True, pin_memory=True)  
test_loader = DataLoader(test_dataset, batch_size=128, shuffle=False, pin_memory=True)
```

```
train_loader = DataLoader(train_dataset, batch_size=128, shuffle=True, pin_memory=True)
test_loader = DataLoader(test_dataset, batch_size=128, shuffle=False, pin_memory=True)
```

b. Enable DataLoader Pinning

Use pin_memory=True in `DataLoader` to speed up data transfer from CPU to GPU.

python

 Copy code

```
train_loader = DataLoader(train_dataset, batch_size=32, shuffle=True, pin_memory=True)
```

has a

gpu
memory