


Recap

02 January 2025 16:49



fashion MNIST

→ pytorch
↓
ANN

60000 images

stage 1
cpu → train

5000 images → 82-83%

stage 2

gpu - train

→ 60000 → 88

The Problem

02 January 2025 16:50

test $\rightarrow$ 88 %
train $\rightarrow$ 98 % $\rightarrow$ 10 %
 $\downarrow$

overfitted

test $\rightarrow$ 88 %
train $\rightarrow$ 93 % $\rightarrow$ 5 %

$\downarrow$
reduce

Solution

02 January 2025 16:50

60000

- 1. Adding more data ✓~~x~~
- 2. Reducing the complexity of NN architecture ~~x~~
- 3. Regularization ✓
- 4. Dropouts ✓
- 5. Data Augmentation → CNN ~~x~~
- 6. Batch Normalization ✓
- 7. Early Stopping ~~x~~


optimize

→ Regularization

→ Dropouts

→ Batch Norm

...



$\frac{\text{Loss} + \text{penalty}}{\text{penalty}}$

train

$\frac{\text{train}}{\text{test}}$

gap
reduce


Dropout

03 January 2025 18:22


$$[p = 0.5] \quad 50\% \\ 0.3$$

1. Applied to the hidden layers
2. Applied after the ReLU activation function
3. Randomly turns off $p\%$ neurons in the hidden layer during each forward pass
4. This has a regularization effect ✓
5. During evaluation dropout is not used

each forward pass




simplify



Batch Norm

03 January 2025 18:40

- Applied to Hidden Layers:
 - Typically applied to the hidden layers of a neural network, but not to the output layer.
- Applied After Linear Layers and Before Activation Functions:
 - Normalizes the output of the preceding layer (e.g., after nn.Linear) and is usually followed by an activation function (e.g., ReLU).
- Normalizes Activations:
 - Computes the mean and variance of the activations within a mini-batch and uses these statistics to normalize the activations.
- Includes Learnable Parameters:
 - Introduces two learnable parameters, gamma (scaling) and beta (shifting), which allow the network to adjust the normalized outputs.
- Improves Training Stability:
 - Reduces internal covariate shift, stabilizing the training process and allowing the use of higher learning rates.
- Regularization Effect:
 - Introduces some regularization because the statistics are computed over a mini-batch, adding noise to the training process.
- Consistent During Evaluation:
 - During evaluation, BatchNorm uses the running mean and variance accumulated during training, rather than recomputing them from the mini-batch.



L2 Regularization

03 January 2025 18:58

Applied to Model Weights:

- Regularization is applied to the weights of the model to penalize large values and encourage smaller, more generalizable weights.

Introduced via Loss Function or Optimizer:

- Adds a penalty term $\lambda \sum w_i^2$ to the loss function in L2 regularization.

$$\text{Loss}_{\text{reg}} = \text{Loss}_{\text{original}} + \lambda \sum w_i^2$$

$$\lambda (w_1^2 + w_2^2 + w_3^2 + w_4^2)$$

- In weight decay, directly modifies the gradient update rule to include λw_i , effectively shrinking weights during training.

$$w \leftarrow w - \eta (\nabla \text{Loss} + \lambda w)$$

weight-decay

Penalizes Large Weights:

- Encourages the network to distribute learning across multiple parameters, avoiding reliance on a few large weights.

Reduces Overfitting:

- Helps the model generalize better to unseen data by discouraging overly complex representations.

Controlled by a Hyperparameter:

- A regularization coefficient (λ , often set via `weight_decay` in optimizers) controls the strength of the penalty. Larger values lead to stronger regularization.

No Effect on Bias Terms:

- Regularization is typically applied only to weights, not biases, as biases don't directly affect model complexity.

Active During Training:

- Regularization affects weight updates only during training. It does not explicitly influence the model during inference.

